

An Introduction to Elang

Joe Armstrong <erlang@gmail.com>

(c) 2018 Joe Armstrong.

Whoomph Software AB

All rights Reserved.

May 16, 2018

Contents

1	To whet your appetite	1
1.1	Patterns	1
1.2	The Multi Server	2
2	Getting started	3
2.1	Starting the shell	3
2.2	Shell commands	4
2.3	Shell commands	4
2.4	Stopping the shell	4
2.5	Recall and edit commands in the shell	5
2.6	Exercise	5
2.7	The shell	5
3	Variables	6
3.1	Variables cannot vary	6
3.2	Back to high school	6
3.3	Single assignment variables	6
4	Pattern matching	7
4.1	Var = value	7

5	Data types	7
5.1	Complex types - Tuples	8
5.2	Creating and unpacking a tuple	8
5.3	Repeated occurrence of variables	8
5.4	Anonymous variables	9
5.5	Complex types - Lists	9
5.6	Lists as recursive structures	9
5.7	Strings	10
5.8	Quiz: If strings are lists...	10
5.9	Complex types - Maps	11
5.10	Exercise	11
5.11	Shell commands	12
6	Compiling and testing	12
6.1	Write the unit test	12
6.2	Correct the code and run	13
6.3	Try again	13
6.4	Anatomy of a function	13
6.5	Useful Trick when developing	14
6.6	Running your program	14
6.7	Or use escript	14
6.8	Code paths etc.	14
6.9	Exercise: During an idle moment...	15
7	Getting help	15
8	Sequential erlang	15
8.1	Review: Types	15
8.2	Functions	16
8.3	Built-In functions (BIFs)	16
8.4	BIFs For type conversion	17
8.5	BIFS for efficiency	17
8.6	BIFS	17
8.7	apply	17
8.8	Funs	18
8.9	Let's add a for loop to Erlang	18
8.10	Funs can return funs	18
8.11	map	19

8.12	Writing our own map	19
8.13	Writing our own map	19
8.14	Accumulators	19
8.15	List comprehensions	20
8.16	Quicksort	20
8.17	Multiple return values	21
8.18	Accumulators	21
8.19	Accumulators	21
8.20	Guards	22
8.21	Single Guard conditions	22
8.22	Full Guard	22
8.23	Guard Sequence	22
8.24	case and if	23
8.25	Booleans	23
8.26	Many functions work with booleans	23
8.27	Tame large tuples with records	24
8.28	Record defaults	24
8.29	Not mentioned	24
8.30	Exercises:	24
9	Exceptions	25
9.1	Convert exceptions into values	25
9.2	Protecting a function call	26
9.3	Protecting a sequence of expressions	26
9.4	Use exceptions for "impossible" situations	27
9.5	Solutions	27
9.6	exit, throw, or error	27
9.7	Defensive programming is painful...	28
9.8	So don't program defensively	28
9.9	Exercise:	29
9.10	Debugging	29
9.11	lib_misc:dump(File, Term)	29
9.12	MACRO: Not Yet Implemented	30
9.13	MACRO: DEBUGGING	30
10	Concurrency	31
10.1	Models of concurrency	31
10.2	Processes or threads?	31

10.3	Concurrency in other languages	31
10.4	Basic ideas	32
10.5	Everything is a process	32
10.6	Processes	32
10.7	Spawn	33
10.8	Send	33
10.9	Receive	33
10.10	Including self() in a message	34
10.11	Spawn send and receive example	34
10.12	Selected or ordered receives	35
10.13	Sending Pids in messages	36
10.14	Counter processes	37
10.15	Another process might reply	38
10.16	Exercise	38
10.17	Data abstraction (1)	39
10.18	Data abstraction (2)	40
10.19	Abstracting Server Functionality	40
10.20	Registered processes.	41
10.21	Client-Server Model	41
10.22	Warnings	42
10.23	Timeouts: receive ... after	42
10.24	Exercise: Write an alarm process	43
10.25	Mutually recursive processes	43
10.26	Spawn process, send the code later	44
10.27	SMP and processes	44
10.28	Summary	44
10.29	Client-Server	45
10.30	Variations: Exit client on server error	46
10.31	Variations: timeout in client	47
10.32	Variations: delegation	47
10.33	Variations: Mobile code	48
10.34	Variations: Mobile code with transactions	48
10.35	Rolling your own	49
11	Links	49
11.1	What is a link?	49
11.2	Messages and signals	50
11.3	Exit messages	51

11.4	Example (1) - monitor	51
11.5	The small print	52
11.6	Crashing and exit propagation	52
11.7	Why do we do this?	53
11.8	Things to think about	54
11.9	Exercise - Keep Alive	55
11.10	The principle of remote error handling	55
11.11	half links (monitor)	55
11.12	Summary	56
12	Sockets Based Distribution	56
12.1	Uses gen_tcp or gen_udp	56
12.2	Client	56
12.3	Server	56
12.4	A sequential server	57
12.5	A parallel server	57
12.6	Packet lengths	57
12.7	Sending Erlang terms	57
12.8	The middle man pattern	58
12.9	Exercise	58
12.10	The Bit Syntax	58
12.11	Tips	59
12.12	Examples	59
13	Distributed Erlang	59
13.1	Distribution primitives	60
13.2	Distributed Erlang	60
13.3	Distributed Erlang	60
13.4	Two nodes, one machine	61
13.5	Shell can connect to remote nodes	61
14	OTP	62
14.1	Structure	62
14.2	Principles	62
14.3	Getting started with applications	63
14.4	behaviors	63
14.5	Case studies	64
14.6	gen_server	64

14.7	gen_supervisor	65
15	Windup	65
15.1	Sequential Erlang	65
15.2	Concurrent Erlang	66
15.3	Fault-tolerant Erlang	66
15.4	Distributed Erlang	66
15.5	Benefits of Erlang	66
15.6	Some Projects to research	66
15.7	And Remember...	67

1 To whet your appetite

By the end of this course I hope you'll be able to *read* Erlang programs - and you'll have started to be able to *write* programs.

To start you off I'm going to show you a small program that by the end of the day you should be able to understand. I'll show it in a moment, but before this we have to talk about patterns.

1.1 Patterns

Patterns are central to Erlang. We'll meet them in many contexts.

In expressions:

Pattern = Value

In function definitions:

```
funcName(Pattern1) -> Actions1;
funcName(Pattern2) -> Actions1;
...
funcName(PatternN) -> ActionsN.
```

In case statements:

```
case Value of
  Pattern1 -> Actions1;
  Pattern2 -> Actions1;
  ...
  PatternN -> ActionsN
end
```

In funs (or closures):

```
F = fun(Pattern1) -> Actions1;
      (Pattern2) -> Actions1;
      ...
      (PatternN) -> ActionsN
end,
```

In message reception:

```
receive
  Pattern1 -> Actions1;
  Pattern2 -> Actions1;
  ...
  PatternN -> ActionsN
end
```

1.2 The Multi Server

```
-module(multi_server).
-compile(export_all).

start() -> spawn(fun() -> loop() end).

loop() ->
  receive
    {_Pid, {email, _From, _Subject, _Text} = Email} ->
      io:format("multi_server:email:~p~n",[Email]),
      {ok, S} = file:open("inbox", [write,append]),
      io:format(S, "~p.~n", [Email]),
      file:close(S);
    {_Pid, {im, From, Text}} ->
      io:format("Msg (~s): ~s~n",[From, Text]);
    {Pid, {get, File}} ->
      io:format("multi_server:get:~s~n",[File]),
      Pid ! {self(), file:read_file(File)};
    Any ->
      io:format("multi server got:~p~n",[Any])
```

```
end,  
loop().
```

This program is:

- An email server - (SMTP becomes an erlang message)
- An instant messaging agent (jabber?, XMPP)
- A web server (or FTP server, or both) (HTTP, ...)

2 Getting started

- Starting the shell
- Entering commands
- Stopping the shell
- Editing commands
- Data types
- Variables
- Pattern matching

2.1 Starting the shell

- Unix-based systems: `$ erl`
- Windows: Programs -> OTP ... -> Erlang

This is what you see

```
$ erl  
Erlang (BEAM) emulator version 5.6 [source] [smp:4] [async-threads:0]  
      [hipe] [kernel-poll:false]  
  
Eshell V5.6  (abort with ^G)  
1>
```


2.2 Shell commands

- Shell is read-eval-print loop
- Commands end with a period (.) followed by whitespace

```
1> 3 * 7.  
21  
2> math:sqrt(2).  
1.41421  
3> "hello". "world".  
"hello"  
4> "world".  
"world"
```

2.3 Shell commands

- Repeated prompt means command is not yet finished

```
5> 123 +  
5> 3456 *  
5> 789.  
2726907  
6>
```

2.4 Stopping the shell

- (ctrl) + C

```
BREAK: (a)bort (c)ontinue (p)roc info (i)nfo (l)oaded  
        (v)ersion (k)ill (D)b-tables (d)istribution  
a
```

- (ctrl) + - immediate exit
- init:stop(). – controlled exit
- erlang:halt() – uncontrolled exit

2.5 Recall and edit commands in the shell

- Arrow keys
- Emacs editor commands
 - Ctrl + B – back
 - Ctrl + F – forward
 - Ctrl + P – previous
 - Ctrl + N – next
- Tab – try to expand module or function names

2.6 Exercise

Try it out. Fire up erlang at a command prompt. Do some basic arithmetic.

```
$ erl
1> 123456789 * 987654321 * 123456789.
15053411111487447638891241
...
```

2.7 The shell

- Major caveat: you cannot define functions in the shell
- Instead, create in separate files and compile in the shell

3 Variables

- Variables start with an upper case letter
- syntax [A..Z] [a..zA..Z0..9] (almost)

```
1> MinutesPerHour = 60.
60
2> HoursPerDay = 24.
24
3> MinutesPerDay = MinutesPerHour * HoursPerDay.
1440
4>
```

3.1 Variables cannot vary

```
4> Count = 1.  
1  
5> Count = Count + 1.  
** exception error: no match of right hand side value 2
```

- These are not like Java or C# variables
- More like variables in math

3.2 Back to high school

$$\begin{aligned}x + 3y &= 15 \\ x - 3y &= 3\end{aligned}$$

$$x = ?, y = ?$$

x and y mean the same thing in all equations.

3.3 Single assignment variables

- As in math, $X = X + 1$ is illegal
- Enables concurrency
- No locks needed since we can't share and update data
- Makes debugging easy (because variable is only set once)
- Lives as long as it is within scope

4 Pattern matching

- `Pattern = Value` matches Pattern against Value
- `X = 10` means "match the pattern X against the value 10"

4.1 Var = value

- if **Var** is a "fresh" (new , unbound) variable, then the pattern is matched by giving **Var** the value **Value**
- if **Var** already has a value, pattern matches if its value equals **Value**, fails otherwise

```
1> X = 10.  
10  
2> X = 10.  % matches  
10  
3> X = 2*4 + 2.  % matches  
10  
4> X = 11.  
** exception error: no match of right hand side value 11
```

5 Data types

Simple data types:

- Integers
 - 1 45 123123 2#101012 8#0723 16#face
- Floats
 - 12.34e-07 3.14159
- Atoms
 - true false hello_world 'funny atom #\$\$!!'
 - atoms are like enumerated types in C
 - (monday, tuesday, wednesday, ... are atoms – these can be used to represent days of the week)

5.1 Complex types - Tuples

- Ordered list of values {food, egg}
- can be nested {person,{firstname,"joe"},{lastname,"armstrong"}}

5.2 Creating and unpacking a tuple

- Constructed using { ... } notation
- Unpacked by pattern matching

```
4> Person = {person, "Joe", "Armstrong"}.
{person, "Joe", "Armstrong"}
5> {Type, First, Last} = Person.
{person, "Joe", "Armstrong"}
6> Type.
person
7> Last.
"Armstrong"
```

5.3 Repeated occurrence of variables

If a variable occurs more than once in a pattern, it must have the same value.

```
1> {X, Y, X} = {1, a, 1}.
{1,a,1}
2> X.
1
3> Y.
a
4> {A, B, A} = {1, a, 2}.
** exception error: no match of right hand side value {1,a,2}
```

5.4 Anonymous variables

`_` is a "don't care" variable. Used as a placeholder in pattern matches.

```
1> {X, _, X} = {1, a, 1}.
{1,a,1}
2> X.
```

5.5 Complex types - Lists

- Sequence of values (can be different types)
- `[]` is the empty list
- `[X1,X2,...,Xn]` is a list with elements `X1`, `X2`, ..., `Xn`
- `[...]` in a pattern deconstructs the list

```
1> X = [1, 2, a, b].  
[1,2,a,b]  
2> [A, B, C, D] = X.  
[1,2,a,b]  
3> A.  
1  
4> D.  
b
```

5.6 Lists as recursive structures

- First element of a list is the *head*, what remains is a list called the *tail*. (Lisp programmers, think `CAR` and `CDR`)
- `[]` is the empty list
- `[H | T]` is the list with head `H` and tail `T` (which must be a list)
- Can be used to create and to pattern match

```
1> X = [1,2,3,4].  
[1,2,3,4]  
2> Y = [abc|X].  
[abc,1,2,3,4]  
3> [X1,_,X2|X4] = Y  
[abc,1,2,3,4]  
4> X3.  
2  
5> X4.  
[3,4]
```

5.7 Strings

- `$a` is short for 97 (The ASCII code for the character `a`)
- `"abc"` is shorthand for `[$a, $b, $c]`
- The shell prints lists of integers as strings (if it can)
- `++` concatenates lists

```
1> Name = "Armstrong".
"Armstrong"
2> "Joe " ++ Name.
"Joe Armstrong"
3> [ 99, 97, 110].
"can"
```

5.8 Quiz: If strings are lists...

```
1> A = "Hi,".
"Hi,"
2> B = "there".
"there"
3> length([A, B]).
?
4> length(A ++ B).
?
5> length([A | B]).
?
```

5.9 Complex types - Maps

```
1> Person = #{first => "Joe", last => "Armstrong"}.
#{first => "Joe",last => "Armstrong"}
2> maps:get(first, Person).
"Joe"
3> maps:get(age, Person).
** exception error: {badkey,age}
   in function maps:get/2
```

```
called as maps:get(age,#{first => "Joe",last => "Armstrong"})
```

```
4> maps:find(age,Person).
error
5> maps:find(first,Person).
{ok,"Joe"}
```

You can pattern match on maps:

```
birthday(#{age := N} = Person) ->
  Person#{arg => N+1}.
```

5.10 Exercise

- Try to predict the result of each of the following *before* typing it into the shell. Use the command `f()`. after each command to *forget* any bindings. `b()` prints all bindings.

```
X = true.
{X,abc} = {123,abc}.
{X,Y,Z} = {222,def,"cat"}.
{X,Y} = {333,ghi,"cat"}
{X,Y,X} = {{abc,12},42,{abc,12}}.
{X,Y,X} = {{abc,12},42,true}.
[H|T] = [1,2,3,4,5].
[H|T] = "cat".
[A,B,C|T] = [a,b,c,d,e,f].
```

5.11 Shell commands

```
> help().
b()          -- display all variable bindings
e(N)         -- repeat the expression in query <N>
f()          -- forget all variable bindings
f(X)         -- forget the binding of variable X
h()          -- history
history(N)   -- set how many previous commands to keep
results(N)   -- set how many previous command results to keep
v(N)         -- use the value of query <N>
```


... many more commands ...

6 Compiling and testing

- Step 1 - Write a unit test
- Step 2 - Write some code
- Step 3 - Iterate 1 and 2

(In most of the exercises I'll give you the unit tests.)

6.1 Write the unit test

```
-module(math3a).  
-compile(export_all).
```

```
test() ->  
    10 = sum([1,2,3,4]).
```

But...

```
1> c(math3a).  
./math3a.erl:5: function sum/1 undefined  
error
```

6.2 Correct the code and run

```
-module(math3b).  
-compile(export_all).
```

```
test() ->  
    10 = sum([1,2,3,4]).
```

```
sum([H|T]) -> H + sum(T).
```

```
1> c(math3b).  
{ok,math3b}  
2> math3b:test().
```

```

** exception error: no function clause matching math3b:sum([ ])
    in function  math3b:sum/1
    in call from math3b:test/0

```

6.3 Try again

```

-module(math3c).
-compile(export_all).

test() ->
    10 = sum([1,2,3,4]),
    -5 = sum([-6,1]),
    0 = sum([]),
    horray.

sum([H|T]) -> H + sum(T);
sum([]) -> 0.

7> c(math3c).
{ok,math3c}
8> math3c:test().
horray

```

6.4 Anatomy of a function

```

%% File name is amod.erl
-module(amod).
%% these can be called from outside the module
-export([foo/1, bar/2]).
%% import this function from lists
-import(lists, [reverse/1]).

foo(L) ->
    L1 = lists:flatten(L),    %% call 'flatten' which must be exported from lists
    reverse(L1).              %% same as lists:reverse(L1)

bar(A, B) -> A + B.

```

6.5 Useful Trick when developing

`-compile(export_all)` exports all functions from this module.

6.6 Running your program

- Compile it `$ erlc mycode`
- Run it `$ erl -noshell -pa /path/to/code -s Mod Func Args ...`

6.7 Or use escript

```
#!/opt/local/bin/escript
```

```
main(List) ->  
    io:format("Args are ~p~n", [List]).
```

Then, from the command line

```
% chmod +x myprog  
% ./myprog 34 67  
Args are ["34","67"]
```

6.8 Code paths etc.

- Erlang autoloads compiled modules using the current code search paths
- For libraries, can set the path in `./.erlang` or `${HOME}/.erlang`
- both are read by `erl` when it starts
- Can contain any `erl` commands

```
code:add_patha("/Users/joe/Desktop/work/2007/jaerlang").  
code:add_patha("/Users/joe/Desktop/work/2007/jaerlang/socket_dist").
```

6.9 Exercise: During an idle moment...

Try adding `io:format("Hello from:~p~n",[file:get_cwd()]).` to either `./.erlang` or `${HOME}/.erlang`.

7 Getting help

Manual pages are not installed by default (but MacPorts does load them for you)

- Unix – `wget http://www.erlang.org/download/otp_doc_man_R11B-5.tar.gz`
 - Unpack so that the `man` directory is unpacked at the root of the Erlang distribution (usually `/usr/local/lib/erlang`).
 - View using `erl -man Modulename`
- Windows – the HTML rendered documentation is included in the distribution

8 Sequential erlang

8.1 Review: Types

Sequential Erlang programs manipulate instances of data types

- integers – `2176537165 1 16#face`
- floats – `1.24234`
- atoms – `monday, tuesday`
- binaries – `<<"3868GAJSJB">>` memory buffers

With two complex data types

- lists – `[X1, X2, ...]`
- tuples – `{X1, X2, ...}`

8.2 Functions

Remember:

```
func(Pattern1) -> Actions1;
func(Pattern2) -> Actions2;
...
func(PatternN) -> ActionsN.
```

Get the punctuation right - SEMI-COLON ... DOT.

```
temp_convert({f, F}) -> {c, 5*(F-32)/9};
temp_convert({c, C}) -> {f, 32 + 9*C/5}.
```

```
lookup(Key, [{Key,Val}|_]) -> {yes, Val};
lookup(Key, [_|T]) -> lookup(Key, T);
lookup(Key, []) -> no.
```

8.3 Built-In functions (BIFs)

- We can't do everything with pattern matching. For example, type conversion BIFs, convert an atom to a list, convert a tuple to a list etc.
- `atom_to_list(Atom)`, `tuple_to_list(Tuple)`, `term_to_binary(term)`, `binary_to_term(Bin)`
- `erl -man erlang` describes all BIFs (book pages 451–460 contains a summary)

8.4 BIFs For type conversion

```
gen_tcp:send(Socket, term_to_binary(Term)), ...
```

```
receive
  {tcp_data, Socket, Bin} ->
    Term = binary_to_term(Bin)
  ...
end.
```

8.5 BIFS for efficiency

```
encode(Term) ->
    B = term_to_binary(Term),
    Md5 = erlang:md5(B),
    <<B/binary,Md5/binary>>.

decode(<<Md5:16/binary,B/binary>>) ->
    Md5 = erlang:md5(B), %% throws an error if wrong value
    binary_to_term(B).
```

8.6 BIFS

- Quirky syntax
- Sometimes `tuple_to_list`, sometimes `erlang:now()`
- In theory the `erlang:` BIFS are implemented in Erlang (not true :-)
- Behave as if they were defined in the module `erlang`

8.7 apply

You can build a call to a function and apply that function at run-time using `apply(Mod, Func, [Arg1, Arg2, ...])`.

```
1> M = list_to_atom("erlang").
erlang
2> F = list_to_atom("tuple_to_list").
tuple_to_list
3> apply(M, F, [{a,b,c}]).
[a,b,c]
4> M:F({a,b,c}).
[a,b,c]
```

8.8 Funs

Funs are “anonymous functions” (also called lambda-expressions, closures).

```
1 > Double = fun(X) -> 2*X end.
#Fun<erl_eval.4.4564646>
2 > Double(2).
4
```

Funs can be arguments to functions.

```
3 > lists:map(Double, [1,2,3,4]).
[2,4,6,8]
```

8.9 Let's add a for loop to Erlang

```
-module(hofs).
-export([for/3]).

for(Max, Max, F) -> [ F(Max) ];
for(I, Max, F)   -> [ F(I) | for(I+1,Max,F) ].

1> c(hofs).
{ok,hofs}
2> hofs:for(1,10,fun(I) -> 2*I end).
[2,4,6,8,10,12,14,16,18,20]
```

8.10 Funs can return funs

```
1> MakeAdder = fun(Inc) -> fun(X) -> X + Inc end end.
#Fun<erl_eval.6.72228031>
2> Add5 = MakeAdder(5).
#Fun<erl_eval.6.72228031>
3> Add5(10).
15
```

8.11 map

- apply the same function to every element of a list. The result is a list the same length as the input list
- `map(F, [X1,X2,...Xn]) -> [F(X1), F(X2), ..., F(Xn)]`

- *Incredibly useful*

```
1 > lists:map(fun(X) -> X * X end, [1,2,3]).
[1,4,9]
```

8.12 Writing our own map

```
my_double([ H | T ]) -> [ 2*H | my_double(T) ];
my_double([]) -> [].
```

8.13 Writing our own map

```
my_double([ H | T ]) -> [ 2*H | my_double(T) ];
my_double([]) -> [].
```

Generalize:

```
my_map(Fun, [ H | T ]) -> [ Fun(H) | my_map(Fun, T) ];
my_map(Fun, []) -> [].
```

```
my_double(L) -> my_map(fun(X) -> 2*X end, L).
```

8.14 Accumulators

Let's write `sum` using an accumulator:

```
sum(L) -> sum(L, 0).
```

```
sum([H|T], Sum) ->
    Sum1 = Sum + H,
    sum(T, Sum1);
sum([], Sum) ->
    Sum.
```

Note we have two different functions, `sum/1` and `sum2/2`. This is a common pattern—the additional parameters are analogous to function-local variables in other languages.

8.15 List comprehensions

```
[ Constructor || Pattern <- List, Predicate, ...]
```

- Come from set theory—the set of all X such that Y , eg $\{ x : x \text{ is even } \}$
- In Erlang, use `||`

```
2> L = lists:seq(1,20).  
[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20]  
3> [ X || X <- L, (X rem 2) =:= 0].  
[2,4,6,8,10,12,14,16,18,20]
```

```
* Allows very compact code  
map(F, L) -> [F(X) || X <- L].
```

8.16 Quicksort

```
-module(qsort).  
-export([sort/1]).  
  
sort([Pivot|T]) ->  
    sort([ X || X <- T, X < Pivot ])  
++ [Pivot] ++  
    sort([ X || X <- T, X >= Pivot ]);  
sort([]) ->  
    [].  
  
1> c(qsort).  
{ok,qsort}  
2> qsort:sort([1,a,3,4,b,c,2]).  
[1,2,3,4,a,b,c]
```

8.17 Multiple return values

- Single return value $X = \text{some_func}(\dots)$
- To get multiple return values use a tuple
 $\{X, Y, Z\} = \text{another_func}(\dots)$

8.18 Accumulators

- Pass extra argument(s) into the function

```
-module(accum).  
-export([evens_and_odds/1]).  
-import(lists, [reverse/1]).  
  
evens_and_odds(L) -> evens_and_odds(L, [], []).  
  
evens_and_odds([H|T], E, 0) when H rem 2 == 0 -> evens_and_odds(T, [H|E], 0);  
evens_and_odds([H|T], E, 0) -> evens_and_odds(T, E, [H|0]);  
evens_and_odds([], E, 0) -> {E, 0}.  
  
1> accum:evens_and_odds([1,2,3,4,5,6,7]).  
{[6,4,2],[7,5,3,1]}
```

8.19 Accumulators

- User reverse to fix order at end

```
-module(accum2).  
-export([evens_and_odds/1]).  
-import(lists, [reverse/1]).  
  
evens_and_odds(L) -> evens_and_odds(L, [], []).  
  
evens_and_odds([H|T], E, 0) when H rem 2 == 0 -> evens_and_odds(T, [H|E], 0);  
evens_and_odds([H|T], E, 0) -> evens_and_odds(T, E, [H|0]);  
evens_and_odds([], E, 0) -> {reverse(E), reverse(0)}.  
  
2> accum:evens_and_odds([1,2,3,4,5,6,7]).  
{[2,4,6],[1,3,5,7]}
```

8.20 Guards

Extend function pattern matching to include additional predicates and comparisons

```
func(A, B) when ...guard... ->
```

For example:

```
func1(A, B) when is_integer(A) ->
  ...
```

8.21 Single Guard conditions

- Type-match predicates (`is_integer(X)`, `is_tuple(X)`, ...)
- Built-in guard functions (`length(X)`, `hd(X)`, ...)
- comparisons, boolean, and arithmetic expressions

```
func1(A, B) when is_integer(A) andalso A + 1 > B ->
  ...
```

8.22 Full Guard

- A comma separated list of single guard conditions
- Matches only if all match

```
func1(A, B) when is_integer(A), is_tuple(B) ->
  ...
```

8.23 Guard Sequence

- A semicolon separated list of full guards
- matches if any match

```
func1(A, B) when is_integer(A), is_tuple(B); is_float(A), is_tuple(B) ->
  ...
```

Or

```
classify(Day) when A == saturday; A == sunday ->
  weekend;
classify(Day) ->
  weekday.
```

8.24 case and if

```
goto_work(Day) ->
  case classify(Day) of
    weekday -> true;
    weekend -> false
  end.
```

```
goto_work(Day) ->
  if
    Day == sunday -> false;
    Day == saturday -> false;
    true -> true
  end.
```

Remember:

```
case Value of
  Pattern1 -> Actions1;
  Pattern2 -> Actions2;
  ...
  PatternN -> ActionsN
end
```

Get punctuation right SEMI-COLONS.

8.25 Booleans

- true and false atoms (not built-in type)
- By convention only, but a strongly recommended convention.

8.26 Many functions work with booleans

- `lists:filter(Fun, L) -> L'`

This takes a list of values $[X_1, X_2, \dots, X_n]$ and produces a new list $[X_i, X_j, \dots]$ of all the vlaues in L for which `Fun(X)` is true.

`lists:partition(Fun, L) -> {T, F}`

Splits the elements X in L into two sub-lists T and F depending on whether `Fun(X)` is true or false.

These functions will not work on things like `on` and `off`.

8.27 Tame large tuples with records

`{person, "joe", "armstrong"}` is fine when there are just a few elements but when we have more than (say) half a dozen elements things get messy.

```
-record(person, {name,firstname,lastname,age,sex,weight,height}).
```

```
X = #person{name="jane", age=32, sex=female} %% creates a new record
```

```
birthday(#person{age=N} = X) ->  
  X#person{age = N+1}
```

Records both construct and pattern match.

8.28 Record defaults

```
-record(window, {width=100, ht=100, x=10, y=10, color=red}).
```

```
X = #window{width=200}, %% other arguments get default values
```

8.29 Not mentioned

- macros
- parse transforms

8.30 Exercises:

Write a program to number the elements in a list. Here's the test case:

```
test() ->  
  [{1,a},{2,b},{3,c}] = number_list([a,b,c]),
```

Write functions `one(List)`, `two(List)`, and `three(List)` that each returns two lists containing the even and odd elements in `List`. Use three different techniques. The test case is:

```
test() ->  
  {[2,4],[1,3,5]} = one([1,2,3,4,5]),  
  {[2,4],[1,3,5]} = two([1,2,3,4,5]),  
  {[2,4],[1,3,5]} = three([1,2,3,4,5]),  
  hooray.
```

9 Exceptions

- Type errors occur when calling a BIF with the wrong type
- Pattern matching errors
- Explicit errors caused by program executing `throw`, `exit`, `error`
- exceptions can be *converted* into values with `catch` and
`try ... catch ... after ... end`

9.1 Convert exceptions into values

```
1> X = atom_to_list(123).
** exited: {badarg,[{erlang,atom_to_list,[1234]}},
               {erl_eval,do_apply,5},
               {erl_eval,expr,5},
               {shell,exprs,6},
               {shell,eval_loop,3}}] **

2> X.
** 1: variable 'X' is unbound **

3> X = (catch atom_to_list(123)).
{'EXIT',{badarg,
  [{erlang,atom_to_list,[1234]}},
  {erl_eval,do_apply,5},
  {erl_eval,expr,5},
  {erl_eval,expr,5},
  {shell,exprs,6},
  {shell,eval_loop,3}}]}}
```

9.2 Protecting a function call

`catch F(X)` returns `F(X)` if no exception is generated while evaluating the function

```
some_function(X, Y) ->
  case (catch something_that_might_crash(X, Y)) of
    {'EXIT', Why} ->
      %% .. handle the error ...;
    NormalReturn ->
      %% regular processing
  end.
```

Evaluating `some_function(X, Y)` will never raise an exception (assuming that we correctly handle the error in the `EXIT` branch of the `case` statement)

9.3 Protecting a sequence of expressions

```
start(Arg1, Arg2) ->
  try
    begin
      ok = start_server(),
      L1 = dothis(Arg1),
      ...
    end
  catch
    exit: Pattern ->
      ...
  after
    ....
end
```

9.4 Use exceptions for "impossible" situations

I once found this code:

```
opcode(load) -> 1;
opcode(store) -> 2;
```

```
opcode(X) ->    %% invalid: what should I do now?
               io:format("bad opcode:~p~n",[X]).
```

What does this return for invalid opcodes?

9.5 Solutions

do something

```
opcode(load) -> 1;
opcode(store) -> 2;
opcode(X) -> exit({ebadOpCode, X}).
```

Solution two: *do nothing, be lazy* (best)

```
opcode(load) -> 1;
opcode(store) -> 2.
```

Program will generate an exception when called with an invalid argument *precisely as in solution two, the only difference is the name of the exception.*

9.6 exit, throw, or error

- use `exit` if this is so serious that the process should die
- use `throw` if you intend to catch the error in the program
- use `error` if you want to behave like a "normal" library error

9.7 Defensive programming is painful...

```
case file:open("config") of
  {ok, T} ->
    case T of
      {port,P,host,H} ->
        case open_socket(Host, Port) of
          {ok, Pid} ->
            ...
          {error, Why} ->
```



```

        ...
        end
    ... ->
    ...
end
...

```

9.8 So don't program defensively

- Use **ONE BIG CATCH AT THE TOP OF YOUR PROGRAM**
- write *everything else* as if it is going to work.
- let your code crash early

```

start(...) ->
  try
    begin
      {ok, T} = file:consult("config"),
      {port,Port,host,Host} = T,
      {ok, Pid} = open_socket(Host, Port)
    end;
  catch
    What:Why ->
      io:format("oooo err: ~p ~p~n",[What,Why])
  end

```

(We'll talk about inter-process errors shortly...)

9.9 Exercise:

The BIF `list_to_integer("1234")` will return 1234, but exits with a bad argument exception if given a bad string. Write a function `list_to_int(String)` that returns `{ok, Int}` or `{error, eBadInt}`.

The unit test is:

```

test() ->
  {ok, 123} = list_to_int("123"),
  {error, eBadInt} = list_to_int("abc"),
  hooray.

```

9.10 Debugging

- There is a debugger—very few people use it
- Errors are pretty easy to find. This is a consequence of single assignment
- Compiler errors are "obvious"
- `lib_misc:dump(File, Term)` for large data structures
- `io:format("...~p....~n", [Var1, ...])`
- `?NYI` macro
- `ifdef(DEBUGGING)`

9.11 `lib_misc:dump(File, Term)`

```
1 > V = epp:parse_file("ex1_bit_syntax.erl","", ""), true.  
true  
2> lib_misc:dump("debug", V).  
Dumping term to debug  
ok
```

`V` is stored in a file called `debug`

```
{ok,  
 [{attribute,1,file,{"ex1_bit_syntax.erl",1}},  
  {attribute,1,module,ex1_bit_syntax},  
  {attribute,2,compile,export_all},  
  {function,  
   31,  
   test,  
   0,  
   [{clause,  
    31,  
    [],  
    [],  
    [{match,
```

```

    32,
    {var,32,'Code'},
    ...

```

9.12 MACRO: Not Yet Implemented

```

-define(NYI(X),(begin
    io:format("*** NYI ~p ~p ~p~n",[?MODULE, ?LINE, X]),
    exit(nyi)
end)).

```

Use like this:

```

...
?NYI({glurk, X, Y}).
...

```

9.13 MACRO: DEBUGGING

```

-define(DEBUGGING, true).
-ifdef(DEBUGGING).
-define(DEBUG(X), X).
-else.
-define(DEBUG(X), void).
-endif.

...
?DEBUG(io:format("Found:~p~n",[Files])),
...

```

10 Concurrency

- The world *is* concurrent
- Programming concurrent activities in a sequential language is *artificially difficult*
- Modern processors are concurrent

10.1 Models of concurrency

- Shared memory (locks, mutexes, coroutines, processes, threads, deadlock, livelock, failure, thread-safe)
- Message passing

10.2 Processes or threads?

- Threads *share* resources (essentially an efficiency hack – beware of premature optimizations)
- Processes do not share things. Processes are the basis of *security* in an operating system
- (Caution) the word “process” is thought of as a synonym for “slow, big, difficult to manage” Erlang processes are *very lightweight*. (Much lighter than a conventional thread)

10.3 Concurrency in other languages

- Can only create very small number of processes (a few thousand at most)
- Heavy weight
- Only supports message passing and not the kind of error handling semantics that Erlang has

10.4 Basic ideas

- If you know the name of a process, you can send it a message
- `Pid ! Message`
- If `Pid` is hidden you cannot send a message to the processes (security)
- Message arrives in a mailbox (like email)
- You never know if a message arrives (if you want to be sure send a message back)

- You can *link* to a process. If a process dies and you linked to it, you will be sent an error signal.
- The error signal is like a message (more later)

10.5 Everything is a process

In the Erlang world we model all non-Erlang things as processes.

We interact with all objects by sending them messages. The behavior of all objects in the external world is inferred from analysing the messages they send to us.

This is the *purest* form of OO (encapsulation, isolation, polymorphic).

10.6 Processes


- `spawn` – create a new process
- `send` – send a message to a process
- `receive` – receive a message

10.7 Spawn


Three syntaxes:

- `spawn(fun foo/0)` – runs `foo/0` in a new process
- `spawn(fun() -> ... end)` – spawns an inline fun
- `spawn(Mod, Func, [Arg1, Arg2, ..., Argn])` performs `apply(Mod, Fun, [Arg1, Arg2, ..., Argn])` in a new process.

In `Pid1` evaluate `Pid2 = spawn(fun() -> ...end)`



Pid1



Pid2

10.8 Send

`Pid ! Msg` sends the message `Msg` to the to the *Mailbox* of a process named `Pid`. `Pid` is the return value of of a previously evaluated `spawn` expression.

10.9 Receive


Receive suspends the process until a message arrives in the mailbox that matches one of the patterns in a `receive` statement. (Rather like a fancy `select` statement)

Syntax:

```
receive
  Pattern1 -> Actions1;
  Pattern2 -> Actions2;
  ...
end
```

10.10 Including `self()` in a message

- `self()` is the pid of the current process
- including `self()` in a message allows the recipient to respond



10.11 Spawn send and receive example

```
-module(counter1).
```

```

-export([start/0]).

start() ->
    Pid = spawn(fun() -> counter(1) end),
    tick(Pid),
    tick(Pid).


tick(Pid) -> Pid ! bump.

counter(N) ->
    io:format("Counter is now ~w~n", [ N ]),
    receive
        bump ->
            io:format("Got bumped~n"),
            counter(N+1)
    end.


1> c(counter1).
{ok,counter1}
2> counter1:start().
Counter is now 1
bump
Got bumped
Counter is now 2
Got bumped
Counter is now 3

```

10.12 Selected or ordered receives




Receive message from A or B whichever comes first



Receive message from A then B

10.13 Sending Pids in messages



10.14 Counter processes

How do we get the value of the counter from the process?

```
-module(counter2).  
-export([start/0, tick/1, read/1]).  
  
start() -> spawn(fun() -> counter(1) end).  
tick(Pid) -> Pid ! bump.  
  
read(Pid) ->  
    Pid ! {self(), read},  
    receive  
        N -> N  
    end.  
  
counter(N) ->  
    receive
```

```

        bump          -> counter(N+1);
        {From, read} -> From ! N, counter(N)
    end.

```

```

1> c(counter2).
{ok,counter2}
2> P = counter2:start().
<0.38.0>
3> counter2:tick(P).
bump
4> counter2:tick(P).
bump
5> counter2:read(P).
3

```

Is anything wrong with this code?

10.15 Another process might reply

“Better” version – include a return value to match on that makes sure the expected process answers.

```

-module(counter3).
-export([start/0, tick/1, read/1]).

start() -> spawn(fun() -> counter(1) end).
tick(Pid) -> Pid ! bump.

read(Pid) ->
    Pid ! {self(), read},
    receive
        {Pid, N} -> N
    end.

counter(N) ->
    receive
        bump          -> counter(N+1);

```

```

        {From, read} -> From ! {self(), N}, counter(N)
    end.

```

10.16 Exercise

`multi_server`

Write `multi_server.erl`. `Pid = multi_server:start()` starts a multi server. It responds to three messages:

- `{Pid, {email, Who, Subject, Text}}` - a request to store email in your inbox
- `{Pid, {im, Who, Text}}` - an instant message - write message to terminal.
- `{Pid, {get, File}}` - a request to fetch a file. The server should reply by sending the file contents (as a message) back to `Pid`.
In all of these `Pid` is the process identifier of the process sending a request to the server.

Hints:

- `{ok, Bin} = file:read_file(File)` reads `File` into a binary.
- `{ok, Stream} = file:open("inbox", [write,append])` opens a file for write append
- `io:format(Stream, "~p.~n", [Term])` writes a term to a file
- `file:close(Stream)` closes a file

10.17 Data abstraction (1)

Did we need to "invent" the functions `tick` and `read`?

```

2> c(counter4).
{ok,counter4}
3> Pid = spawn(fun() -> counter4:counter(0) end).
<0.42.0>
4> Pid ! bump.

```

```

bump
5> Pid ! bump.
bump
6> counter4:rpc(Pid, read).
2

-module(counter4).
-export([counter/1, rpc/2]).

rpc(Pid, What) ->
    Pid ! {self(), What},
    receive
        {Pid, N} -> N
    end.

counter(N) ->
    receive
        bump          -> counter(N+1);
        {From, read} -> From ! {self(), N}, counter(N)
    end.

```

10.18 Data abstraction (2)

- Are we hiding the wrong thing? Should we expose the protocol?
- PROS: Hides the messaging structure. Makes "module resuability" easy. We know how to define APIs.
- CONS: More code. We don't know how to describe protocols. Protocols do not exist in Erlang (or in any other language for that matter) as first class objects.

10.19 Abstracting Server Functionality

```

-module(abs_server).
-export([start/1, cast/2, rpc/2]).

start(Data) ->

```

```

        spawn(fun() -> loop(Data) end).

cast(Pid, F) -> Pid ! {cast, F}.

rpc(Pid, F) ->
    Pid ! {rpc,self(), F},
    receive
        {Pid, Reply} -> Reply
    end.

loop(State) ->
    receive
        {cast, F} -> loop(F(State));
        {rpc, From, F} ->
            {Reply, State1} = F(State),
            From ! {self(), Reply},
            loop(State1)
    end.

-module(counter5).

-export([start/0, tick/1, read/1]).

start() -> abs_server:start(0).

tick(Pid) -> abs_server:cast(Pid, fun(X) -> X+1 end).
read(Pid) -> abs_server:rpc(Pid, fun(X) -> {X, X} end).

> Pid1 = counter5:start().
<0.36.0>
> counter5:tick(Pid1).
{cast,#Fun<counter5.0.20406117>}
> counter5:tick(Pid1).
{cast,#Fun<counter5.0.20406117>}
> counter5:read(Pid1)
2

```

10.20 Registered processes.

Any program that wants to send a message to `Pid` must have `Pid` in a local variable. This must be communicated to all functions that want to send a messages to `Pid`

This is secure (but inconvenient)

- `register` – registers a global name with a `Pid`
- `unregister` * After `register(name, Pid)`, `name ! Term` sends a message to `Pid`
- `whereis(Name)` returns `Pid` or `undefined`

10.21 Client-Server Model

```
-module(reg_counter1).
-export([start/0, tick/0, read/0]).

start() ->
    register(counter, spawn(fun() -> counter(1) end)).

tick() -> counter ! bump.

read() ->
    counter ! {self(), read},
    receive
        {counter, N} -> N
    end.

counter(N) ->
    receive
        bump -> counter(N+1);
        {From, read} -> From ! {counter, N}, counter(N)
    end.
```

10.22 Warnings

If two processes do `register` with the *same name at the same time* one will succeed and one will fail.

Suppose we evaluate the following in two parallel processes:

```
make_global(Name, F) ->
    register(Name, spawn(F)).
```


One will fail the other will succeed.

- Can you write a process `make_global(Name, Fun)` that is correct using only `register` and `whereis`?
- Think about it

10.23 Timeouts: receive ... after

```
receive
    Pattern1 -> ...
    Pattern2 -> ...
after
    TimeInMilliseconds ->
        Actions
end.
```

Timeouts



10.24 Exercise: Write an alarm process

- `alarm(Pid, Msg, Time)` sends the message `Msg` to `Pid` after `Time` milliseconds.
- Use it to write "Hello, World!" to your console every two seconds

10.25 Mutually recursive processes

Process one needs to know `Pid2` and processes two needs to know `Pid1`

```
Pid1 = spawn(fun() -> start(F1) end),  
Pid2 = spawn(fun() -> start(F2) end),  
Pid2 ! Pid1,  
Pid1 ! Pid2,  
...  
start(F) ->  
    receive
```



```

        Pid -> F(Pid)
    end

```

10.26 Spawn process, send the code later

```

Pid = spawn(fun() -> wait() end),
...
Pid ! {do, F}
...

wait() ->
    receive
        {do, F} ->
            F()
    end

```

10.27 SMP and processes

- Erlang can take advantage of multiple processors and cores
- Enable using `-smp s n + options` (only if compiled in)

10.28 Summary

All concurrent behavior is programmed with


- `spawn`
- `send`
- `receive`

Instead of *blocks*, *mutexes*, *threads*, *processes*, *synchronised methods*,
 ...

10.29 Client-Server


The Most Important Concurrency Pattern

Client Server Model



Protocol

- Protocol



```
-module(s1).  
-export([start/0, f1/2, f2/3, ..., stop/0]).  
  
start() -> register(s1, spawn(fun() -> loop(State1))).  
  
%% interface routines  
f1(A, B) -> rpc(s1, {do_f1,A,B}).
```

```

f2(...) -> rpc(s1, ...).

rpc(Q) ->
    s1 ! {self(), Q},
    receive
        {s1, Reply} -> Reply
    end.

loop(State) ->
    receive
        {From, {do_f1, A, B}} ->
            ...
            FFrom ! {s1, Reply},
            loop(State1);
        ...
    end.

```

10.30 Variations: Exit client on server error

Exit the client if the RPC causes an error in the server

```

loop(State) ->
    receive
        ..
        {From, Q} ->
            try
                {Reply, State1} = f(Q, State),
                From ! {Name, ok, Reply},
                loop(State1);
            catch
                exit:Why ->
                    From ! {Name, die, Why},
                    loop(State)
            end
    end

rpc(Name, Q) ->
    Name ! {self(), Q},

```

```

receive
  {Name, ok, Reply} -> Reply;
  {Name, die, Why} -> exit(Why)
end.

```

10.31 Variations: timeout in client

```

rpc(Name, Q) ->
  Name ! {self(), Q},
  receive
    {Name, ok, Reply} -> Reply;
    {Name, die, Why} -> exit(Why)
  after
    Time ->
      exit(timeout)
  end.
...

```

10.32 Variations: delegation

```

loop() ->
  receive
    {From, Q} ->
      ...
      Other ! {do, Q, replyAs, Name, replyTo, From}

```

Delegated

```

loop() ->
  receive
    {do, Q, replyAs, Name, replyTo, From} ->
      ...
      From ! {Name, Reply}

```

Responder

```

loop() ->
  receive
    {From, Q} ->
      ...
      From ! {Name, Reply}

```

Original

10.33 Variations: Mobile code

```

loop(State) ->
  ...
  {From, Func} ->
    {Reply, State1} = Func(State),
    From ! {self(), Reply},
    loop(State1);
  ...

```

10.34 Variations: Mobile code with transactions

```

loop(State) ->
  ...
  {From, Func} ->
    try Func(State) of
      {Reply, State1} ->
        From ! {self(), Reply},
        loop(State1);
    catch
      _:_ ->
        exit(From, evilLaugh),
        loop(State)
    end
  ...

```


10.35 Rolling your own

- As you can see there are many variations of the client-server theme.
- No two ways of writing the client-server loop do exactly the same thing. These programs can be made to do exactly what you want.
- If you understand this then you can easily design your own middleware system

11 Links

- `link` – link to a process
- `unlink` – remove the link
- `process_flag(trap_exit, true)` – receive a signal as a message

11.1 What is a link?




- `links` define error propagation paths

- If A is linked to B then: *If A fails then B will be notified If B fails then A will be notified*
- links are not "stacked" - calling `link` ten times has the same effect as calling it once.
- links propagate at the speed of messages - they are not instantaneous (important in distributed systems)

11.2 Messages and signals

- Messages are sent with `send`
- Signals are sent when processes die. Signals are not messages
- If A is linked to B, C, ... then B,C,... will be sent an exit signal if A dies
- All processes die when they receive signals, unless they are set to trap exits
- To trap exits a process evaluates `process_flag(trap_exit, true)`




11.3 Exit messages

- {'EXIT', Pid, Why} can be received if you are trapping exits
- If you are *not* trapping exits and Why is not the atom `normal` you die

11.4 Example (1) - monitor

`monitor(Pid)` observer Pid and print a message if Pid dies.





```
monitor(Pid) ->
  link(Pid),
  process_flag(trap_exit, true),
  receive
    {'EXIT',Pid,Why} ->
      io:format("Process died Reason:~p~n",[Pid,Why])
  end.
```


11.5 The small print


- You can fake an exit message by calling `exit(Pid1, Why)` – If you are `Pid2` then `Pid1` sees `{'EXIT',Pid2, Why}` and assumes you have died - but you are really alive. *Do not abuse.*
- exit reason `kill` is unstoppable.

11.6 Crashing and exit propagation






The process marked with a *double ring* is an error trapping process.



- Exits propagate through linked processes.
- If any of A..F crash all processes crash
- System processes stop the propagation of errors.
- Process D does not crash

11.7 Why do we do this?

We can build the system in layers to make it fault tolerant.



11.8 Things to think about

- If you do `spawn` followed by `link` the process might die very quickly (ie, before getting to the `link` statement)
- Solution: `spawn_link` is like `spawn` followed by `link` only the two are performed atomically
- Need to make sure `trap_exit` is evaluated **before** we try to catch errors – think about synchronization here

11.9 Exercise - Keep Alive

Write a function `make_global(Name, Fun/0)` that starts a "global" process named `Name`. If `Name` dies for any reason, restart it. Record the times and reasons for starting and restarting in an error log.

Useful stuff:

- `time() -> {Hour, Minute, Second}`
- `error_logger:format(String, List)` writes to the error log

11.10 The principle of remote error handling

- To make a fault-tolerant system you need two (separated) computers.
- Let one process do the work
- Let some other process fix the error
- You can think of exit signals as uncaught exceptions that escape from the process and propagate to some other process
- Processes that fail should die-early
- We can make very reliable systems this way
- We can test on one node and deploy on multiple nodes (everything works the same way)

11.11 half links (monitor)

- `erlang:monitor(process, Pid)`
- asymmetric
- countable (if a process is monitored N times it must be demonitored N times before the monitor is released - unlike link/unlink)

11.12 Summary

- `link` - create a link
- `unlink` - release the link
- `process_flag(trap_exit, true)` – trap exits
- `{'EXIT',Pid,Why}` then message sent when a process dies

All of this works in distributed Erlang. These mechanism are orthogonal to the `spawn`, `send` and `receive`

Now we've seen all that you need to make a powerful fault-tolerant system.

12 Sockets Based Distribution

12.1 Uses `gen_tcp` or `gen_udp`

Book chapter 14

12.2 Client

```
{ok, Socket} = gen_tcp:connect(Host, Port, [Options]),
ok = gen_tcp:send(Socket, Data),
...
receive
    {tcp, Socket, Data} ->
        %% do something with the data
        ...
    {tcp_closed, Socket} ->
        %% take care of this ...
```

12.3 Server

A "single-shot" server (accepts one connection)

```
start_server() ->
    {ok, Listen} = gen_tcp:listen(Port, [Options]),
    {ok, Socket} = gen_tcp:accept(Listen),
    %% We use same calls as the client to read and write
    %% the socket
```

12.4 A sequential server

```
{ok, Listen} = gen_tcp:listen(Port, ...)
seq_loop(Listen).

seq_loop(Listen) ->
    {ok, Socket} = gen_tcp:accept(Listen),
    loop(Socket),
    seq_loop(Listen).
```

12.5 A parallel server

Now turn this into a parallel server

```
{ok, Listen} = gen_tcp:listen(Port, ...)
spawn(fun() -> par_loop(Listen) end).

par_loop(Listen) ->
    {ok, Socket} = gen_tcp:accept(Listen),
    spawn(fun() -> par_loop(Listen) end),
    loop(Socket).
```

12.6 Packet lengths

A 4 byte length header is automatically added/removed by the system when calling `gen_tcp:send` and messages are assembled to the correct length before `{tcp, Sock, Data}` messages are sent to the controlling process

```
gen_tcp:connect(Host, Port, [..., {packet,4}, ...])
gen_tcp:listen(Port, [..., {packet, 4}, ...])
```

12.7 Sending Erlang terms

```
gen_tcp:send(Socket, term_to_binary(Term)) ...
```

```
receive
    {tcp, Socket, Data} ->
        Term = binary_to_term(Data),
        ...
```

12.8 The middle man pattern

```
loop(Pid, Socket) ->
    receive
        {Pid, Msg} ->
            gen_tcp:send(Socket, term_to_binary(Msg)),
            loop(Pid, Socket);

        {tcp, Socket, Data} ->
            Pid ! binary_to_term(Data),
            loop(Pid, Socket);

        {'EXIT', Pid} ->
            gen_tcp:close(Socket);

        {tcp_closed, Socket} ->
            exit(Pid, socket_closed)
    end.
```

12.9 Exercise

Do this at home!

`multi_server` over sockets

Write `server.erl` and `client.erl`.

Use a middle man to connect a client on one machine to a server on another machine. Use `middle_man.erl` to connect the client to a socket on one machine and on the other machine to connect a process owning a socket to an instance of `multi_server`.

Hint: Book page 248 (and these notes) have the parallel server pattern. The middle man pattern is described in the course notes and on page 404 of the book in the description of `lib_chan_mm`

12.10 The Bit Syntax

Pack/Match bit strings from binaries

```
1> Red = 2, Green=61, Blue=20.
20
```

```

%% Red takes 5 bits, Green 6, and blue 5
2> Bin = <<Red:5, Green:6, Blue:5>>.
<<23,180>>
3> io:format("~8.2.0B ~8.2.0b~n", binary_to_list(Bin)).
00010111 10110100
%% 00010 111101 10100

```

Size must be a multiple of 8 bits

12.11 Tips

- Can handle different endian words, big, little, ...
- complex syntax
- experiment in shell then cut and paste into program

12.12 Examples

- Book page 86 - find MP3 headers (for syncing with SHOUTCAST server)
- Too numerous to mention
- Fun just now. Decode/Encode AMF (Actionscript binary protocol)
- Useful for writing assemblers, protocol convertors etc.

13 Distributed Erlang

- Distributed Erlang – several **nodes** belong to the same system. **spawn** has an extra argument **spawn(Node, Mod, Fun, Args)**
- SMP Erlang – symmetric multiprocessing. Two or more identical processors are connected to a single shared main memory.
- Socket Distribution – Not really distributed at all

13.1 Distribution primitives

Adds three new primitives and a load of libraries

- `spawn(Node, Mod, Func, Args)` – *everything* works as before regarding `links` exit messages an so on.
- `alive(Node)` – tell the system you are alive

The main libraries

- `rcp` – doing remote procedure calls
- `global` – global operation over all nodes

13.2 Distributed Erlang

- Cookie based security
- Distribution over TCP (can be over secure sockets, but installation is more complex)
- Only suitable in a cluster
- Used in enterprise software behind a firewall

13.3 Distributed Erlang

The most common architectural pattern is to write traffic handling as non distributed applications on a single node. We use `mnesia` and the error loggers running replicated on multiple nodes.

`mnesia` the (Erlang database) is can be configured to provide table replication over multiple nodes.

Fault tolerant applications usually use replicated pairs of nodes. One node does the work the other is a hot standby.

13.4 Two nodes, one machine

```
$ erl -sname one          %% in one terminal window
(one@joe-armstrongs-computer)1>

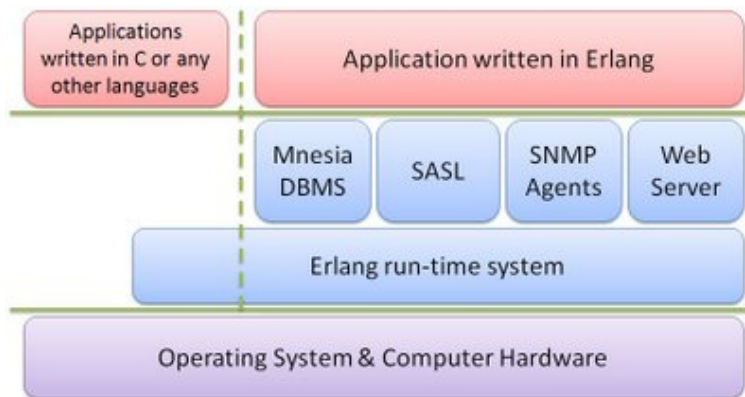
$ erl -sname two          %% in another window
> (two@joe-armstrongs-computer)1> net_adm:ping('two@joe-armstrongs-computer').
pong
> (two@joe-armstrongs-computer)1> net_adm:ping('one@joe-armstrongs-computer').
pong
(two@joe-armstrongs-computer)2> node().
'two@joe-armstrongs-computer'
(two@joe-armstrongs-computer)3> rpc:call('one@joe-armstrongs-computer',
                                         erlang, node, []).
'one@joe-armstrongs-computer'
```

13.5 Shell can connect to remote nodes

```
$erl -name one
..
(one@joe-armstrongs-computer.local)1> ^G
User switch command
--> j
    1* {shell,start,[init]}
--> r 'server+doris.myerl.home.net'
--> j
    1 {shell,start,[init]}
    2* {'server@doris.myerl.home.net',shell,start,[]}
--> c 2
Eshell V5.5.5 (abort with ^G)
(server+doris.myerl.home.net)1> node().
'server@doris.myerl.home.net' ^G
User switch command
--> c 1
(one+joe-armstrongs-computer.local)1> node().
'one@joe-armstrongs-computer.local'
```

14 OTP

OTP Architecture



- Open Telecomms Platform
- Maintained by product group inside Ericsson
- Releases 2-3 times/year
- "<http://www.erlang.org/>":<http://www.erlang.org/>
- "<http://www.erlang.org/doc.html>":<http://www.erlang.org/doc.html>

14.1 Structure

Generic structure

```
/appname/ebin %% beam code
      /src    %% erlang source
      /priv   %% everything else
```

14.2 Principles

- "http://www.erlang.org/doc/design_principles/part_frame.html":http://www.erlang.org/doc/design_principles/part_frame.html
- Overview

- Client-server `gen_server`
- Finite State machines `gen_fsm`
- Event handling `gen_event`
- Supervisor `gen_sup`
- Applications
- Releases
- Application upgrade

14.3 Getting started with applications

- Read book chapters 16 and 18 *especially* +16.1
- Read PhD thesis "<http://www.sics.se/joe/thesis>":<http://www.sics.se/joe/thesis>
- Read *design principles* "http://www.erlang.org/doc/design_principles/part_frame.html"
- Forthcoming O'Reilly book

14.4 behaviors

- The OTP name for "design patterns"
- Callback modules for client-servers, supervision trees etc.
- Encapsulates "best practice" from many individual projects
- 3'rd generation - ie the third rewrite of basic servers
- Used in practice

14.5 Case studies

See Armstrong PhD thesis.

- AXD301 – `gen_server` (122), `gen_event` (36), `supervisor` (20), `gen_fsm` (10), `application` (6).
- Nortel Networks – `gen_server` (56), `supervisor` (19), `application` (15), `gen_event` (9), `rpc_server` (2), `gen_fsm` (1), `supervisor_bridge` (1).


`gen_server` is the most used behavior

14.6 `gen_server`

- A generic client-server model
- `gen_server:start_link(Name, Mod, InitArgs, Opts)`
- `Mod:handle_call(...)` gets called for all RPCs
- `Mod:handle_cast(...)` gets called for all casts
- ...
- `Mod:code_change(...)` gets called when you want to change the code

More later...

14.7 gen_supervisor



- Supervision trees
- Hierarchical tree of workers and supervisors
- Supervisors monitor the workers
- If a worker fails this is noticed by the supervisor
- Supervisors can start, restart and kill workers
- 1:1 and 1:N supervision

15 Windup

15.1 Sequential Erlang

- Simple functional language
- Lists, tuples, atoms, bignums, floats, ...
- Function selection is by pattern matching

- Data selection is by pattern matching
- Variables are immutable

15.2 Concurrent Erlang

- Adds `spawn`, `send` and `receive` to sequential Erlang.
- `register` [`unregister`] can be used to associate a name with a process

15.3 Fault-tolerant Erlang

- `catch .. throw` and `try ... catch ... end` added to sequential Erlang
- `link`, `process_flag(trap_exit, true)` added to concurrent Erlang

15.4 Distributed Erlang

- Add `+spawn(Node, Mod, Func, Args)` to concurrent Erlang
- Or use explicit term passing over sockets

15.5 Benefits of Erlang

- Multi-core ready
- Processes in the language (not OS)
- Designed for fault-tolerant distributed programming
- Battle tested

15.6 Some Projects to research

- AXD301 (Ericsson)
- Kreditor (kreditor.se)
- SimpleDB (Amazon)
- CouchDB (text db)

- MociWeb (Mochimedia)
- Ejabberd (jabber server)
- ErlyWeb (Erlang on Rails :-)

15.7 And Remember...

Make It Fun!